


## Deep Learning - Udacity



Just as with neural networks, we create a CNN in Keras by first creating a `Sequential` model.

We add layers to the network by using the `.add()` method.

Copy and paste the following code into a Python executable named `conv-dims.py`:

```
from keras.models import Sequential
from keras.layers import Conv2D

model = Sequential()
model.add(Conv2D(filters=16, kernel_size=2, strides=2, padding='valid',
                  activation='relu', input_shape=(200, 200, 1)))
model.summary()
```

We will not train this CNN; instead, we'll use the executable to study how the dimensionality of the convolutional layer changes, as a function of the supplied arguments.

Run `python path/to/conv-dims.py` and look at the output. It should appear as follows:

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 100, 100, 16)	80
=====		
Total params: 80		
Trainable params: 80		
Non-trainable params: 0		

Do the dimensions of the convolutional layer line up with your expectations?

Feel free to change the values assigned to the arguments (`filters`, `kernel_size`, etc) in your `conv-dims.py` file.

Take note of how the **number of parameters** in the convolutional layer changes. This corresponds to the value under `Param #` in the printed output. In the figure above, the convolutional layer has 80 parameters.

Also notice how the **shape** of the convolutional layer changes. This corresponds to the value under `Output Shape` in the printed output. In the figure above, `None` corresponds to the batch size, and the convolutional layer has a height of 100, width of 100, and depth of 16.

### Formula: Number of Parameters in a Convolutional Layer

The number of parameters in a convolutional layer depends on the supplied values of `filters`, `kernel_size`, and `input_shape`. Let's define a few variables:

- $K$  - the number of filters in the convolutional layer
- $F$  - the height and width of the convolutional filters
- $D_{in}$  - the depth of the previous layer

Notice that  $K = \text{filters}$ , and  $F = \text{kernel\_size}$ . Likewise,  $D_{in}$  is the last value in the `input_shape` tuple.

Since there are  $F \times F \times D_{in}$  weights per filter, and the convolutional layer is composed of  $K$  filters, the total number of weights in the convolutional layer is  $K \times F \times F \times D_{in}$ . Since there is one bias term per filter, the convolutional layer has  $K$  biases. Thus, the **number of parameters** in the convolutional layer is given by  $K \times F \times F \times D_{in} + K$ .

### Formula: Shape of a Convolutional Layer

The shape of a convolutional layer depends on the supplied values of `kernel_size`, `input_shape`, `padding`, and `stride`. Let's define a few variables:

- $K$  - the number of filters in the convolutional layer
- $F$  - the height and width of the convolutional filters
- $S$  - the stride of the convolution
- $H_{in}$  - the height of the previous layer
- $W_{in}$  - the width of the previous layer

Notice that  $K = \text{filters}$ ,  $F = \text{kernel\_size}$ , and  $S = \text{stride}$ . Likewise,  $H_{in}$  and  $W_{in}$  are the first and second value of the `input_shape` tuple, respectively.

The **depth** of the convolutional layer will always equal the number of filters  $K$ .

If `padding = 'same'`, then the spatial dimensions of the convolutional layer are the following:

- **height** =  $\text{ceil}(\text{float}(H_{in}) / \text{float}(S))$
- **width** =  $\text{ceil}(\text{float}(W_{in}) / \text{float}(S))$

If `padding = 'valid'`, then the spatial dimensions of the convolutional layer are the following:

- **height** =  $\text{ceil}(\text{float}(H_{in} - F + 1) / \text{float}(S))$
- **width** =  $\text{ceil}(\text{float}(W_{in} - F + 1) / \text{float}(S))$

Please change the `conv-dims.py` file, so that it appears as follows:

```
from keras.models import Sequential
from keras.layers import Conv2D

model = Sequential()
model.add(Conv2D(filters=32, kernel_size=3, strides=2, padding='same',
                 activation='relu', input_shape=(128, 128, 3)))
model.summary()
```

Run `python path/to/conv-dims.py`, and use the output to answer the questions below.

### Question 1 of 3

How many parameters does the convolutional layer have?

- 
- 
- 
- 

### Question 2 of 3

What is the depth of the convolutional layer?

- 
- 
- 
- 

### Question 3 of 3

What is the width of the convolutional layer?

- 
- 
- 
-